

# Zagatti\_HW05

May 31, 2022

## 1 Homework 05 - Sebastiano Zagatti

```
[2]: !pip install pyro-ppl

import os
import sys
import json
import numpy as np
import torch
torch.set_default_dtype(torch.float64)

import sklearn
from sklearn.datasets import make_circles
from sklearn.gaussian_process import GaussianProcessClassifier
import sklearn.gaussian_process as gp_sklearn

import pyro
import pyro.distributions as dist
from pyro.infer import MCMC, HMC, NUTS
from pyro.infer import SVI, Trace_ELBO, TraceEnum_ELBO
from pyro.contrib.autoguide import AutoDiagonalNormal
from pyro.optim import Adam
import pyro.contrib.gp as gp

import matplotlib.pyplot as plt
import seaborn as sns
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>  
Requirement already satisfied: pyro-ppl in /usr/local/lib/python3.7/dist-packages (1.8.1)  
Requirement already satisfied: numpy>=1.7 in /usr/local/lib/python3.7/dist-packages (from pyro-ppl) (1.21.6)  
Requirement already satisfied: pyro-api>=0.1.1 in /usr/local/lib/python3.7/dist-packages (from pyro-ppl) (0.1.2)  
Requirement already satisfied: torch>=1.11.0 in /usr/local/lib/python3.7/dist-packages (from pyro-ppl) (1.11.0+cu113)  
Requirement already satisfied: opt-einsum>=2.3.2 in

```
/usr/local/lib/python3.7/dist-packages (from pyro-ppl) (3.3.0)
Requirement already satisfied: tqdm>=4.36 in /usr/local/lib/python3.7/dist-
packages (from pyro-ppl) (4.64.0)
Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.7/dist-packages (from torch>=1.11.0->pyro-ppl) (4.2.0)
```

Let's consider a binary classification problem on Circles dataset. The input is two-dimensional and the response is binary (0,1).

We observe 1000 points  $x$  from this dataset and their labels  $y$ :


```
[3]: x, y = make_circles(n_samples=1000, factor=0.3, noise=0.05, random_state=0)
x = torch.from_numpy(x)
y = torch.from_numpy(y).double()

def scatterplot(x, y):
    colors = np.array(['0', '1'])
    sns.scatterplot(x[:, 0], x[:, 1], hue=colors[y.int()])

scatterplot(x, y)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variables as keyword args: x, y. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
```

FutureWarning



## 1.1 scikit-learn: GaussianProcessClassifier

1. GaussianProcessClassifier from scikit-learn library approximates the non-Gaussian posterior by a Gaussian using Laplace approximation. Define an RBF kernel `gp_sklern.kernels.RBF` with `lengthscale` parameter = 1 and fit a Gaussian Process classifier to the observed data  $(x, y)$ .

```
[4]: RBF = gp_sklern.kernels.RBF(length_scale = 1.0)
GPC = GaussianProcessClassifier(kernel = RBF, random_state = 0)

fitted_GPC = GPC.fit(x, y)
```

2. Use `plot_sklern_predictions` function defined below to plot the posterior predictive mean function over a finite grid of points. You should pass as inputs the learned GP classifier `sklern_gp_classifier`, the observed points  $x$  and their labels  $y$ .


```
[5]: def meshgrid(x, n, eps=0.1):
    x0, x1 = np.meshgrid(np.linspace(x[:, 0].min()-eps, x[:, 0].max()+eps, n),
                        np.linspace(x[:, 1].min()-eps, x[:, 1].max()+eps, n))
    x_grid = np.stack([x0.ravel(), x1.ravel()], axis=-1)
    return x0, x1, x_grid

def plot_sklern_predictions(sklern_gp_classifier, x, y):
    x0, x1, x_grid = meshgrid(x, 30)
    preds = sklern_gp_classifier.predict_proba(x_grid)
    preds_0 = preds[:,0].reshape(x0.shape)
    preds_1 = preds[:,1].reshape(x0.shape)
    plt.figure(figsize=(10,6))
    plt.contourf(x0, x1, preds_0, 101, cmap=plt.get_cmap('bwr'), vmin=0, vmax=1)
    plt.contourf(x0, x1, preds_1, 101, cmap=plt.get_cmap('bwr'), vmin=0, vmax=1)
    plt.title(f'Posterior Mean')
    plt.xticks([]); plt.yticks([])
    plt.colorbar()
    scatterplot(x, y)
```

```
[6]: plot_sklern_predictions(GPC, x, y)
```

/usr/local/lib/python3.7/dist-packages/seaborn/\_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



## 1.2 scikit-learn: GaussianProcessClassifier

Consider the following generative model:

$$\begin{aligned}
 y_n | p_n &\sim \text{Bernoulli}(p_n) & n = 1, \dots, N \\
 \mu &\sim \mathcal{N}(0, 1) \\
 \sigma, l &\sim \text{LogNormal}(0, 1) \\
 \text{logit}(\mathbf{p}) | \mu, \sigma, l &\sim \mathcal{GP}(\mu, K_{\sigma, l}(x_n))
 \end{aligned}$$

We model the binary response variable with a Bernoulli likelihood. The logit of the probability is a Gaussian Process with predictors  $x_n$  and kernel matrix  $K_{\sigma, l}$ , parametrized by variance  $\sigma$  and lengthscale  $l$ .

We want to solve this binary classification problem by means of HMC inference, so we need to reparametrize the multivariate Gaussian  $\mathcal{GP}(\mu, K_{\sigma, l}(x_n))$  in order to ensure computational efficiency. Specifically, we model the logit probability as:

$$\text{logit}(\mathbf{p}) = \mu \cdot \mathbf{1}_N + \eta \cdot L$$

where  $L$  is the Cholesky factor of  $K_{\sigma, l}$  and  $\eta_n \sim \mathcal{N}(0, 1)$ . This relationship is implemented by the `get_logits` function below.

```
[7]: def get_logits(x, mu, sigma, l, eta):
      kernel = gp.kernels.RBF(input_dim=2, variance=torch.tensor(sigma),
      ↪ lengthscale = torch.tensor(1))
      K = kernel.forward(x, x) + torch.eye(x.shape[0]) * 1e-6
      L = K.cholesky()
      return mu+torch.mv(L,eta)
```

3. Write a pyro model `gp_classifier(x,y)` that implements the reparametrized generative model, using `get_logits` function and `pyro.plate` on independent observations  $y_n$ .

```
[8]: def gp_classifier(x, y):
      mu = pyro.sample("mu", dist.Normal(0, 1))
      sigma = pyro.sample("sigma", dist.LogNormal(0, 1))
      l = pyro.sample("l", dist.LogNormal(0, 1))
      with pyro.plate('i', len(x)):
          eta = pyro.sample("eta", dist.Normal(0, 1))
          logit_p = get_logits(x, mu, sigma, l, eta)
          y = pyro.sample("y", dist.Bernoulli(torch.sigmoid(logit_p)), obs = y)
```

4. Use pyro NUTS on the `gp_classifier` model to infer the posterior distribution of its parameters. Set `num_samples=10` and `warmup_steps=50`. Then extract the posterior samples using `pyro.get_samples()` and print the keys of this dictionary using `.keys()` method.

```
[10]: kernel = NUTS(gp_classifier)
      mcmc = MCMC(kernel, warmup_steps=50, num_samples=10, num_chains = 3)
      mcmc.run(x, y)
```

```
/usr/local/lib/python3.7/dist-packages/pyro/infer/mcmc/api.py:500: UserWarning:
num_chains=3 is more than available_cpu=1. Chains will be drawn sequentially.
```

```
num_chains, available_cpu
Warmup: 0%|          | 0/60 [00:00, ?it/s]/usr/local/lib/python3.7/dist-
packages/ipykernel_launcher.py:2: UserWarning: To copy construct from a tensor,
it is recommended to use sourceTensor.clone().detach() or
sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).
```

```
Sample [0]: 100%|      | 60/60 [57:22, 57.37s/it, step size=6.84e-02, acc.
prob=0.204]
Sample [1]: 100%|      | 60/60 [54:28, 54.47s/it, step size=3.56e-02, acc.
prob=0.412]
Sample [2]: 100%|      | 60/60 [58:09, 58.15s/it, step size=2.22e-03, acc.
prob=0.665]
```

```
[11]: samples = mcmc.get_samples()
      samples.keys()
```

```
[11]: dict_keys(['eta', 'l', 'mu', 'sigma'])
```

The `posterior_predictive` function below outputs the prediction corresponding to the  $i$ -th sample

from the posterior distribution. `plot_pyro_predictions` calls this method to compute the average prediction on each input point and plots the posterior predictive mean function over a finite grid of points.

```
[12]: def posterior_predictive(samples, i, x, x_grid):
    kernel = gp.kernels.RBF(input_dim=2, variance=samples['sigma'][i],
                           lengthscale=samples['l'][i])
    N_grid = x_grid.shape[0]

    y = get_logits(x, samples['mu'][i], samples['sigma'][i],
                  samples['l'][i], samples['eta'][i])

    with torch.no_grad():
        gpr = gp.models.GPRegression(x, y, kernel=kernel)
        mean, cov = gpr(x_grid, full_cov=True)

        yhat = dist.MultivariateNormal(mean, cov + torch.eye(N_grid) * 1e-6).
        sample()
        return yhat.sigmoid().numpy()

def plot_pyro_predictions(posterior_samples, x):
    n_samples = posterior_samples['sigma'].shape[0]
    x0, x1, x_grid = meshgrid(x, 30)
    x_grid = torch.from_numpy(x_grid)
    preds = np.stack([posterior_predictive(posterior_samples, i, x, x_grid) for i in
        range(n_samples)])
    plt.figure(figsize=np.array([10, 6]))
    plt.contourf(x0, x1, preds.mean(0).reshape(x0.shape), 101,
                cmap=plt.get_cmap('bwr'), vmin=0, vmax=1)
    plt.title(f'Posterior Mean')
    plt.xticks([]); plt.yticks([])
    plt.colorbar()
    scatterplot(x, y)
```

5. Pass the learned posterior samples obtained from NUTS inference and the set of training points  $x$  to `plot_pyro_predictions` and plot the posterior predictive mean.

```
[13]: plot_pyro_predictions(samples, x)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: UserWarning: To
copy construct from a tensor, it is recommended to use
sourceTensor.clone().detach() or
sourceTensor.clone().detach().requires_grad_(True), rather than
torch.tensor(sourceTensor).
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning:
Pass the following variables as keyword args: x, y. From version 0.12, the only
valid positional argument will be `data`, and passing other arguments without an
```

explicit keyword will result in an error or misinterpretation.  
FutureWarning

